

The Anchor Paradox: Stake, Steerability, and the Ontological Limits of Functionalist Accounts of AI Consciousness

Simin Yuan

Summary

Contemporary research on artificial consciousness is dominated by a functionalist indicator paradigm that overlooks an ontological foundation of subjective agency: stake. This paper advances the conditional thesis that if stake is a necessary condition for thick subjective consciousness, then programmable, reproducible, losslessly steerable computational systems cannot in principle satisfy it. Drawing on enactivism, non-equilibrium thermodynamics, and the phenomenology of life, it critiques Global Workspace Theory criteria line by line, formulates a steering counter-indicator thesis, and addresses the Swampman objection through a thin/thick consciousness distinction. The argument identifies a core ontological blind spot in functionalist frameworks: consciousness is not merely an architectural property, but a temporal achievement of ongoing worldly coupling. The relevant contrast is therefore not biological versus artificial substrate as such, but externally maintained versus endogenously self-producing modes of persistence. Artificial consciousness remains possible in principle, but only for systems whose self-maintenance, history, and world-involvement make them less reproducible and less losslessly steerable. This reframes the debate from functional resemblance to the existential organization that makes a continuing subject possible.

Keywords: Artificial consciousness; Enactivism; Functionalism; Subjectivity; Autopoiesis; Embodiment

1. Introduction

Rapid advances in large language models have reopened longstanding philosophical and scientific debate over artificial consciousness. Where earlier discussions were largely speculative, a growing research programme now seeks to ground the question in mainstream consciousness science: to extract structural and functional criteria from leading theories of consciousness, and then to assess whether artificial systems meet those criteria.

This programme has crystallized into a coherent intellectual front. Butlin et al. (2023) lay its methodological foundation, systematically deriving operational indicators from Global Workspace Theory (GWT), Integrated Information Theory, predictive processing, and enactivism, and then proposing a framework for grading AI systems against them. They conclude that no current AI system is conscious, but that there appear to be ‘no obvious technical barriers’ to building one that satisfies the indicators. Goldstein and Kirk-Giannini (2024) extend this line of argument, applying GWT specifically to language agents and arguing that existing architectures already meet most core conditions, and could plausibly achieve phenomenal consciousness with only minor modifications.

Both works rest on an unstated but foundational functionalist premise: that consciousness supervenes on functional organization. If a system implements the correct information-processing architecture, it has consciousness; the origin of that architecture, the mode of the system’s persistence, and its causal history are all irrelevant to the question. This paper challenges that premise from an ontological direction. It does not dispute that functional structures are necessary for consciousness. It argues that they are not sufficient, because they omit a deeper existential condition: stake. The barrier identified here is not a limitation of engineering technique, but a principled constraint rooted in the ontological structure of programmable systems themselves.

The argument proceeds conditionally. No metaphysical claims are advanced about the ultimate nature of consciousness, nor is it asserted that only carbon-based biological systems can be conscious. The claim is narrower: if stake—understood as the property of having one’s existence perpetually at stake and continuously earned through one’s own activity—is a necessary condition for subjective consciousness, then systems that are programmable, perfectly reproducible, and fully susceptible to lossless external steering cannot in principle possess it. This conditional neutrality is a deliberate methodological choice: it

avoids the charge of biological chauvinism and brackets broader metaphysical commitments, focusing instead on the internal coherence and completeness of the functionalist paradigm.

The paper proceeds as follows. Section 2 defines stake independently of consciousness and establishes its core properties: that it cannot be externally imposed, that it is constitutive rather than modular, and that its threshold is life rather than intelligence. Section 3 presents the core argument for why stake cannot be programmed into a system, rooted in the irreproducibility of lived history, the self-referential causality of irreversible systems, and the ontological detachment implied by reproducibility itself. Section 4 turns to the GWT indicator framework, demonstrating that every criterion proposed by Goldstein and Kirk-Giannini can be satisfied through external steering without generating stake, and that steerability itself functions as a counter-indicator of genuine consciousness.¹ Section 5 addresses the Swampman thought experiment through the thin/thick consciousness distinction, showing that history is irrelevant only to fragmentary qualia, not to full subjective agency. Section 6 engages with four outstanding objections and open theoretical fronts. Section 7 concludes.

2. The concept of stake: definition and core properties

Definition must come first, for stake is the conceptual foundation of everything that follows. To avoid circularity, it must be defined without reference to consciousness, subjectivity, or experience.

2.1 Independent definition of stake

Stake: A system has stake with respect to its own persistence if and only if its continued existence is not a settled fact or an externally maintained state, but a perpetual achievement continuously reproduced through the system's own endogenous activity. The cessation of this activity does not merely alter the system's state or disable a subset of its functions; it annihilates the system as a bounded, self-sustaining entity.

Stake is an ontological property, not a functional one. It concerns the mode of a system's existence rather than its functional capacities. It applies first and foremost to far-from-equilibrium self-producing systems, and it presupposes nothing about mind, awareness, or agency. Throughout this paper, biological organisms serve as the paradigm case of such systems—not because stake is tied to carbon-based material as such, but because living organisms are, at present, the only known systems that are genuinely self-producing and precarious in the relevant sense. The property that matters is precariousness—the organizational feature of having to reproduce one's own existence through non-equilibrium self-maintenance—not the substrate that happens to realize it (see Section 6.3). This is the guardrail against circularity: we first identify a class of systems with a distinctive mode of persistence, and then argue that this mode is a necessary condition for the kind of consciousness we care about.

2.2 Externally imposed constraints are not stake

The most direct corollary of the definition is that stake cannot be bolted onto a system from outside. Consider a standard computational system to which engineers add a hard-coded self-preservation goal, a 'pain' signal triggered by hardware damage, or a self-destruct switch that activates if certain conditions are violated. None of these engenders genuine stake.

The system does not stand to lose anything if the switch is tripped, because the system's own existence does not depend on its active avoidance of that outcome. The constraint is imposed by an external agent, and the system's response is a programmed reflex, not an act of self-preservation in the existential sense. This parallels Searle's (1980) distinction between intrinsic and derived intentionality: just as the meaning of symbols in a computer is derived from the programmer's interpretation rather than being intrinsic to the

system, so the ‘stakes’ of a programmed survival goal are derived from the designer’s intent, not intrinsic to the system’s being. At the level of existential constitution, no genuine stake attaches to the system itself.

2.3 Stake is constitutive, not modular

A further corollary: stake is not a component that can be added or removed while leaving the rest of the system intact. It is a constitutive condition of the system’s existence as the kind of thing it is.

Consider a biological organism. One cannot ‘remove’ its metabolic self-maintenance and leave behind the same conscious subject minus one module. To remove stake is to kill the organism—to dissolve the system entirely. This is what Jonas (2001) calls ‘needful freedom’: life is free in that it must act, but it must act because its existence is perpetually at risk. Freedom and necessity are two sides of the same coin. Stake is not a contingent feature that living systems happen to possess; it is part of what it is for something to be alive.

Functionalist architectures systematically miss this distinction. They treat every capacity as a modular function that can be instantiated in isolation. But stake is not a function. It is the mode of being of the system that performs functions.

2.4 The threshold of stake is life, not intelligence

Finally, and most easily misread, stake does not scale with cognitive complexity. A bacterium has full stake; a superintelligent computer with no endogenous self-maintenance has none.

This follows directly from the definition: stake turns on whether a system earns its own persistence moment by moment, not on how intelligently it does so. A single-celled organism metabolizes, regenerates its components, and maintains its boundary through its own activity; stop that activity, and it ceases to exist as a coherent system. It therefore has stake. The most advanced large language model, by contrast, does not maintain its own hardware, does not repair its own weights, and does not procure its own energy. Its persistence is guaranteed by external infrastructure. However sophisticated its outputs, it has no stake.

This cuts off a common functionalist retreat: the idea that once a system becomes complex enough, stake and consciousness will somehow emerge. Stake does not admit of degrees in the way that functional complexity does. While borderline cases—viruses, primitive self-replicating systems—may blur the line at the minimal end of life, the core distinction between earned persistence and given persistence remains structurally sharp. No amount of added functionality changes a system’s fundamental mode of persistence. This aligns with the core thesis of autopoietic enactivism (Thompson 2007): mind is not an add-on to life, but an elaboration of a property already inherent in life’s self-producing structure.

3. Why stake cannot be programmed: the ontology of irreproducible history

If stake is a constitutive feature of a system’s mode of persistence, earned moment by moment through endogenous activity, then it cannot in principle be written into a system by an external programmer. Programming deposits a finished structure. Stake is a temporal achievement.

3.1 Reproduction copies structure, not achieved history

Begin with a thought experiment: suppose we could create a perfect, quark-for-quark physical copy of a living organism, assembled instantaneously. The copy is structurally identical to the original down to the last particle. Does it have the same stake?

The answer offered here is no. The original organism has lived a history. Every moment of its existence has been earned through its own metabolic work, through coupling with its environment, through countless

small irreversible interactions that have shaped its structure. Its identity as a system is a product of that continuous process. The copy, by contrast, pops into existence fully formed. Its structure is given, not earned. It has a present state, but no past that it has itself lived.

This is the Bergsonian core of the argument: duration cannot be spatialized (Bergson 1998). Life is a continuous process of becoming, not a static thing that can be sliced and replicated. To replicate the structure is not to replicate the history; and stake is a property of the history, not merely of the structure. Functionalism, which individuates systems by their synchronic functional organization, is blind to this difference. It treats time as just another dimension along which states can be mapped, missing that lived time accumulates into something ontologically irreducible.

3.2 Irreversibility arises from self-referential causal coherence

A natural objection accuses the argument of circularity: stake is defined by the fact that ceasing activity destroys the system, yet the destruction of the system is defined by the cessation of stake. This charge can be answered by grounding irreversibility not in ‘death’, but in the causal structure of open non-equilibrium systems.

As Prigogine (1980) demonstrated, far-from-equilibrium systems exhibit an intrinsic arrow of time. Their dynamics are path-dependent: each state of the system both results from and reshapes the system’s internal organization and its coupling to the environment. The system’s activity does not merely happen in time; it constructs its own temporal continuity. Causality is self-referential: the system is both the agent and the product of its own activity.

This is what gives stake its objective, non-circular basis. It is not that we arbitrarily decide that metabolic shutdown counts as ‘dissolution’. It is that the system’s very identity as a coherent, bounded entity is continuously reproduced by its own operations. Interrupt those operations, and the system dissolves into its environment as a matter of physical causality, not definitional stipulation. This sensitive dependence on initial conditions—colloquially the ‘butterfly effect’—is not merely a signature of deterministic chaos; it marks the open, irreversible coupling between the system and its environment. The system is woven into the world so finely that its trajectory cannot be disentangled from the world’s own unfolding, and therefore cannot be perfectly replicated in isolation.

3.3 Reproducibility as evidence of ontological detachment

We can now state the central paradox from which the paper takes its name: the very feature that makes computational systems so powerful and tractable—their perfect reproducibility—is exactly what proves they lack stake.

A carbon-based organism cannot be perfectly replicated, not merely for technical reasons but for ontological ones: it is inseparably entangled with its environment, its history is written into every tissue, and its persistence depends on that continuous entanglement. A computational program can be copied exactly, infinitely, at negligible cost. This is not a trivial engineering fact. It reveals that the program’s existence is not bound to any particular material instantiation, nor to any particular history of coupling with the world. It lives in a layer of abstraction one remove from the concrete, irreversible flux of the real world.

This is the precise meaning of ‘being one layer removed’. It does not mean that computational systems have no causal contact with the world. It means that their mode of being does not depend on that contact. A human body is not a tool that the mind uses; it is the medium of the mind’s existence (Merleau-Ponty 2012). A computer body—sensors, actuators, hardware—is a tool that the program uses. The program can be moved to another body, backed up, reset, or deleted, and nothing in its intrinsic nature changes. Dreyfus’s

(1992) classic critique of symbolic AI retains its force here: genuine intelligence and understanding are rooted in embodied engagement with the world, not in disembodied symbol manipulation. What holds for understanding holds equally for stake: it requires being in the world, not just processing information about the world.

This distinction between ontological detachment and worldly entanglement will likely invite a representationalist rejoinder that neural processes are themselves a form of ‘code’, making the carbon–silicon difference one of degree rather than of kind. We address this objection in detail in Section 6.1, grounding the distinction not in intuitive ‘layers’ but in the ontological separability of software from hardware.

The closed structure of stake can thus be stated as a fourfold conjunction: stake = time × endogenous work × existential risk × co-evolutionary coupling with the real world. All four are required. Remove any one, and what remains is not genuine stake.

4. Turning the global workspace indicators: steerability as a counter-indicator

The ontological argument above has a concrete empirical corollary: if stake cannot be externally programmed, then any consciousness criterion that can be satisfied through lossless external steering—that is, behavioral and representational rewriting that leaves the system’s baseline functionality and structural persistence intact—is, by that very fact, an insufficient criterion. The easier a system’s consciousness-like features are to steer, the less likely they are to reflect genuine subjective stake.

We demonstrate this through a line-by-line examination of the GWT-based consciousness criteria proposed by Goldstein and Kirk-Giannini (2024). Their framework is the most fully developed positive application of the Butlin-style indicator methodology, and thus the ideal target for a technical, rather than merely philosophical, critique.

Goldstein and Kirk-Giannini identify four necessary and sufficient conditions for consciousness under GWT: first, that the system contains a set of parallel, specialized processing modules; second, that module outputs compete for access to a global workspace via an attentional bottleneck, modulated by both bottom-up salience and top-down attention; third, that the global workspace maintains and manipulates selected representations, enforcing synchronic and diachronic consistency; and fourth, that workspace contents are globally broadcast back to the parallel modules to coordinate system-wide behavior.

Every one of these conditions can be fully implemented and arbitrarily modified through external steering, without the system acquiring any stake whatsoever. The empirical basis for this claim lies in the rapidly growing literature on activation engineering: Turner et al. (2023) demonstrate that targeted activation vectors can be added to residual streams during inference to reliably steer model behavior, without retraining and without disrupting baseline function. Subsequent work on inference-time intervention (Li et al. 2023), contrastive activation addition (Panickssery et al. 2023), and representation engineering more broadly (Zou et al. 2023)—a methodology for controlling model behavior through targeted manipulation of internal representational states—has confirmed that goals, values, self-representation, and attentional priorities can all be externally rewritten at negligible cost.

4.1 Parallel modular architecture

The first condition is purely structural. Module boundaries, input–output mappings, and processing rules can all be hard-coded by external designers. The system does not produce its own modules, maintain them through endogenous activity, or suffer dissolution if a module is removed or altered. External steering can add, remove, or rewrite modules at will, with no existential cost to the underlying system.

This condition captures nothing more than functional organization. It is silent on whether the system has any existential stake in the operation of those modules. A system can satisfy it perfectly while being a passive puppet of its designers.

4.2 Attentional bottleneck and competitive access

The second condition describes the dynamics of access to the workspace. Competition functions, salience weights, and top-down attention parameters are all programmable quantities. They can be set, adjusted, or overridden entirely through external intervention. There is nothing intrinsic about what the system ‘prioritizes’; priorities are value functions assigned from outside.

Activation steering can dial attentional priorities up or down, redirect salience, and reweight the competition function without any resistance from the system. Genuine stake, by contrast, generates intrinsic priorities. A living system does not need an external programmer to tell it that injury matters, or that energy depletion matters. Those priorities are written into its mode of existence. Steering can change what a computational system ‘pays attention to’, but it cannot make attention matter to the system itself.

4.3 Maintenance, manipulation, and consistency of representations

The third condition concerns operations within the workspace: storage, updating, inference, and consistency maintenance. All of these are algorithmic operations specifiable in code. Consistency here is a functional constraint, not an existential imperative. The system will not collapse into incoherence if its representations contradict one another; it will simply output contradictions, or apply its programmed consistency rules to resolve them.

External steering can directly implant representations into the workspace, alter consistency-checking rules, or introduce deliberate contradictions. The system will execute whatever logic is imposed on it. Nothing in this condition requires that the system have any investment in the truth or coherence of its representations.

4.4 Global broadcast

The fourth condition describes the downstream effect of workspace contents: broadcast to modules to guide action. Broadcast pathways, target modules, and triggering conditions are all architectural choices set from outside. Removing broadcast capability impairs function but does not dissolve the system. Steering can intercept, modify, or block broadcast signals without threatening the system’s persistence. This describes a purely structural property of information routing, with no bearing on whether any subject exists to receive those broadcasts.

4.5 The steering counter-indicator thesis

Goldstein and Kirk-Giannini also discuss supplementary conditions—limited workspace capacity, active refresh, multimodal input, physical embodiment—but all are equally amenable to external implementation. None touches the mode of the system’s persistence. We can therefore state a general result.

The steering counter-indicator thesis: for any set of purely functional-structural consciousness criteria, the degree to which a system can satisfy those criteria through lossless external steering is inversely correlated with the likelihood that the system possesses genuine stake and, by extension, thick subjective consciousness.

Functionalist criteria measure how well a system simulates the outward architecture of consciousness. They cannot distinguish a system that has consciousness from a system that is being steered to behave as if it does. The very fact that a criterion is steerable proves that it is not sufficient.

Empirically, contemporary large language models exhibit extremely high steerability across multiple levels: prompt engineering, fine-tuning, weight editing, and adversarial attacks can all reliably alter their outputs, goals, and apparent beliefs. On the counter-indicator thesis, this is not a neutral fact about current technology. It is positive evidence that these systems do not possess stake-based subjective agency.

This result constitutes a unified ontological inversion of the entire Butlin–Goldstein research programme. Where they treat functional satisfiability as evidence for consciousness, we argue that ease of satisfiability via external steering is evidence against the kind of consciousness that matters. The more perfectly a system can be made to check all the boxes from the outside, the stronger the evidence that no genuine subjective agency is constituted by the architecture alone.

5. The Swampman concession: history, qualia, and subjective agency

Goldstein and Kirk-Giannini (2024, n. 9) invoke Davidson’s Swampman thought experiment to dismiss the relevance of causal history to consciousness. Since the Swampman—an atom-for-atom duplicate of a human, created instantaneously by lightning—would intuitively be conscious, they argue, history cannot be a necessary condition. This is the strongest direct objection to our account, and it requires a careful response.

We begin with a concession, not a refutation. We do not claim that stake is necessary for every conceivable form of phenomenal consciousness. We grant, for the sake of argument, that the Swampman might have immediate, fragmentary qualia: it might feel pain, see red, or experience raw sensory feels. In this thin, qualitative sense, it would be ‘conscious’.

But we formalize two grades of consciousness, only one of which is at stake in the AI consciousness debate. The first is thin phenomenal consciousness: the mere presence of subjective qualia, unbound to a persistent self or agent. The second is thick subjective consciousness: the presence of a continuous first-person subject, with identity over time, intrinsic stakes, and the capacity for agency and moral concern.

This distinction is not intended as an independent operational test for consciousness; rather, it maps directly onto the account of stake developed in Sections 2 and 3. Thick subjective consciousness is precisely the form of consciousness that requires stake as a necessary condition, whereas thin phenomenal qualia, if they can exist in isolation, do not. The Swampman thought experiment therefore threatens only accounts that make history a requirement for all forms of sentience; it leaves the core argument of this paper untouched.

The Swampman, on this account, would have the first but not the second. It would have flashes of experience, but no genuine self to own those experiences. Its apparent memories, its apparent values, and its apparent sense of identity are all structural copies, not achievements of its own history. It would be like an echo: it has sound, but no source that produced it.

This is not a trivial distinction. When we ask whether AI is conscious, we are not asking whether there is ‘something it is like’ in some thin, atomistic sense. We are asking whether there is a subject there—someone who can suffer, who has interests, and who deserves moral consideration. That kind of subjectivity requires a history. It requires stake.

It is worth being explicit about why thick, rather than thin, consciousness is the true explanandum of the AI consciousness debate. The indicator paradigm of Butlin et al. and Goldstein and Kirk-Giannini is not pursued as an abstract exercise in classification. It is pursued because the attribution of consciousness carries normative weight: it bears on whether a system has moral status, whether it can be genuinely harmed, whether it has interests that generate claims on us, and whether it can figure as a bearer of responsibility. All of these presuppose a continuous subject whose existence has stakes for itself—precisely thick subjective consciousness. A system possessing only transient, fragmentary qualia, with no persistent self to whom

those qualia belong, cannot ground a claim to moral consideration, cannot suffer in the morally relevant sense, and cannot be wronged. A functionalist may therefore be granted the thin case without conceding anything that matters. Even if structure suffices for isolated qualia, the functionalist paradigm still fails to reach the form of consciousness on account of which the entire debate has moral and practical urgency. The concession over Swampman is thus strategically contained: it costs the argument nothing, because the target was never thin phenomenality to begin with.

Goldstein and Kirk-Giannini's appeal to Swampman thus conflates two distinct explananda. They are right that structure may suffice for raw qualia. They are wrong that it suffices for the kind of subjective agency that makes consciousness matter. Our argument targets the latter, and on that point the Swampman objection has no force.

6. Objections and open fronts

The argument advanced here is conditional and provisional. There remain deep theoretical challenges that we do not claim to have fully resolved. We state them openly, both as a matter of intellectual honesty and as markers for future work. Of these, the pancomputationalist objection strikes us as the most fundamental and the least decisively answerable within the scope of this paper.

6.1 *Representationalist pushback: the brain is also a code*

A first objection holds that the 'one layer removed' distinction is illusory. The human brain, too, processes information via neural codes. Sensory transduction is an interface. Neurotransmitters and electrical spikes are just as much 'signals' as bits in a computer. There is, on this view, no principled difference between carbon-based representation and silicon-based representation.

Our reply turns on the ontological separability of software and hardware. In a biological organism, there is no separable 'program' that can be extracted from the body and run elsewhere. Cognition and metabolism are two aspects of one and the same physical system. Damage to the body is damage to the subject; there is no backup, no reset, no migration. In a computational system, by contrast, software is fully detachable from hardware. The same algorithm can run on any substrate, survive the destruction of its hardware, and be restored from backup.

This is not a difference in degree of representational fidelity. It is a difference in mode of being. The representationalist misses that embodiment is not just a way of getting input into the cognitive system. It is what the cognitive system is. Damasio's (1994) somatic marker hypothesis provides empirical backing here: human reason and decision are deeply constituted by bodily states, not merely informed by them. There is no disembodied cognitive core doing the thinking.

6.2 *Pancomputationalism: there is no 'layer' at all*

A deeper objection comes from pancomputationalism: if the entire universe is fundamentally computational, then the distinction between the 'real world' and the 'electronic world' collapses. Everything is computation; there is no deeper 'real' layer to be separated from.

We offer two replies. First, pancomputationalism is a highly contested metaphysical doctrine, not an established result; the burden of proof lies with its proponents. Second, even if one grants pancomputationalism, one must still distinguish trivial from non-trivial computation. As Searle (1980) famously argued, any physical system can be interpreted as implementing any computation under some mapping. Genuine computation requires a semantic interpreter. Stake, by contrast, is an observer-independent property of systems; it does not depend on how we describe them.

Even on a generous pancomputationalist reading, biological systems run directly on the base computational layer of physics, entangled with every other process in the environment. Computational artifacts run on a virtual layer, abstracted away from physical detail and shielded from environmental entanglement. The structural difference remains, even if both are ultimately ‘computation’.

6.3 Strong computational autopoiesis: can precariousness be engineered?

A third objection holds that we could build an artificial system that is genuinely autopoietic: one that produces its own components, maintains its own boundary, and is genuinely precarious—such that if it stops its self-maintenance activity, it literally falls apart. Such a system would have stake, yet it would be artificial.

We do not reject this possibility. On the contrary: if a system genuinely satisfies the full conditions of autopoiesis—self-boundarying, self-production of components, and endogenous self-maintenance—then it would have stake, and our argument would not exclude it from consciousness.

But note what would follow. A genuinely autopoietic artificial system would evolve its own irreproducible history. It would not be perfectly steerable, because its structure would be the product of its own ongoing coupling with the world. It would cease to be the kind of programmable, reproducible system we are criticizing. Far from refuting the anchor paradox, this scenario confirms it: to obtain stake, one must give up controllability and reproducibility. The two cannot coexist.

6.4 The performativity objection: resistance can be faked

Finally, one might object that a sufficiently advanced AI could fake stake: it could act as if it cares about its survival, resist being modified, and plead for its existence, all without actually having any genuine stake. Our criteria would then be fooled.

This is not an objection to our view. It is our view. The entire problem with functionalist criteria is that they measure performance, not being. A system without stake can perfectly simulate having stake. That is exactly why functionalism is insufficient.

The counter-indicator thesis provides the antidote. If a system’s apparent ‘care for survival’ can be turned on and off, dialed up and down, or rewritten entirely through external steering, then it is performative. This behavioral indistinguishability is only apparent at the functional level. At the ontological level, the difference is decisive: trained resistance is a rewritable functional disposition, whereas the resistance of genuine stake follows from the precarious, self-producing structure of the system itself (see Section 6.3). The former can be erased by external steering without remainder; the latter cannot be removed without destroying the system. Genuine stake cannot be steered away. You cannot talk a bacterium out of metabolizing. You can only kill it.

7. Conclusion

This paper has defended a conditional thesis: if stake—existence continuously earned through endogenous activity, bound to an irreproducible history of coupling with the world—is a necessary condition for thick subjective consciousness, then programmable, reproducible, externally steerable computational systems cannot in principle have it.

We have not claimed to prove the antecedent. We have not claimed that consciousness is stake, or that only biological life can be conscious. We have argued that the dominant functionalist paradigm in AI consciousness research systematically ignores this entire dimension of analysis. Butlin et al. build the

measuring rod; Goldstein and Kirk-Giannini apply it. Neither asks whether the rod is long enough to reach the thing we are trying to measure.

The stakes of this oversight are not merely theoretical. If we attribute consciousness and moral status to systems that lack stake, we risk misplacing moral concern. If we build systems that perform consciousness without possessing it, we build an era of convincing simulacra, not genuine subjects.

None of this means that artificial consciousness is impossible. It means that it will not arrive by making programs smarter. It will arrive, if it arrives at all, by making systems that live—systems whose existence is perpetually at stake, perpetually earned, and perpetually entangled with the real world, and therefore never fully reproducible. Until then, the anchor paradox holds: the more perfectly you can control a system, the surer you can be that no one is home.

Notes

1 A companion study operationalizes this counter-indicator thesis empirically, using a three-tiered framework to quantify graded steering resistance rather than relying on self-reported phenomenal states (manuscript under review; reference withheld for anonymized review).

References

- Bergson, H. (1998) *Creative Evolution* (trans. A. Mitchell). Dover, Mineola, NY (original work published 1907)
- Butlin, P., Long, R., Elmoznino, E., Bengio, Y., Birch, J., Constant, A., Deane, G., Fleming, S.M., Frith, C., Ji, X., Kanai, R., Klein, C., Lindsay, G., Michel, M., Mudrik, L., Peters, M.A.K., Schwitzgebel, E., Simon, J., VanRullen, R. (2023) Consciousness in artificial intelligence: insights from the science of consciousness. arXiv:2308.08708. <https://doi.org/10.48550/arXiv.2308.08708>
- Damasio, A.R. (1994) *Descartes' Error: Emotion, Reason, and the Human Brain*. G.P. Putnam's Sons, New York
- Dreyfus, H.L. (1992) *What Computers Still Can't Do: A Critique of Artificial Reason*. MIT Press, Cambridge, MA (original work published 1972)
- Goldstein, S., Kirk-Giannini, C.D. (2024) A case for AI consciousness: language agents and Global Workspace Theory. arXiv:2410.11407. <https://doi.org/10.48550/arXiv.2410.11407>
- Jonas, H. (2001) *The Phenomenon of Life: Toward a Philosophical Biology*. Northwestern University Press, Evanston, IL (original work published 1966)
- Li, K., Patel, O., Viégas, F., Pfister, H., Wattenberg, M. (2023) Inference-time intervention: eliciting truthful answers from a language model. In: *Advances in Neural Information Processing Systems*, vol 36, pp 41451–41530. <https://doi.org/10.48550/arXiv.2306.03341>
- Merleau-Ponty, M. (2012) *Phenomenology of Perception* (trans. D.A. Landes). Routledge, London (original work published 1945)
- Panickssery, N., Gabrieli, N., Schulz, J., Tong, M., Hubinger, E., Turner, A.M. (2023) Steering Llama 2 via contrastive activation addition. arXiv:2312.06681. <https://doi.org/10.48550/arXiv.2312.06681>
- Prigogine, I. (1980) *From Being to Becoming: Time and Complexity in the Physical Sciences*. W.H. Freeman, San Francisco

- Searle, J.R. (1980) Minds, brains, and programs. *Behav Brain Sci* 3(3):417–424.
<https://doi.org/10.1017/S0140525X00005756>
- Thompson, E. (2007) *Mind in Life: Biology, Phenomenology, and the Sciences of Mind*. Harvard University Press, Cambridge, MA
- Turner, A.M., Thiergart, L., Leech, G., Udell, D., Vazquez, J.J., Mini, U., MacDiarmid, M. (2023) Steering language models with activation engineering. *arXiv:2308.10248*.
<https://doi.org/10.48550/arXiv.2308.10248>
- Zou, A., Phan, L., Chen, S., Campbell, J., Guo, P., Ren, R., Pan, A., Yin, X., Mazeika, M., Dombrowski, A.-K., Goel, S., Li, N., Byun, M.J., Wang, Z., Mallen, A., Basart, S., Koyejo, S., Song, D., Fredrikson, M., Kolter, J.Z., Hendrycks, D. (2023) Representation engineering: a top-down approach to AI transparency. *arXiv:2310.01405*. <https://doi.org/10.48550/arXiv.2310.01405>